



```
import pandas as pd
```

```
data = pd.read_csv('/content/weight-height.csv')
```

```
#checking first 10 elements of dataset
```

```
data.head(10)
```

| | Gender | Height | Weight |  |
|---|--------|-----------|------------|---|
| 0 | Male | 73.847017 | 241.893563 |  |
| 1 | Male | 68.781904 | 162.310473 | |
| 2 | Male | 74.110105 | 212.740856 | |
| 3 | Male | 71.730978 | 220.042470 | |
| 4 | Male | 69.881796 | 206.349801 | |
| 5 | Male | 67.253016 | 152.212156 | |
| 6 | Male | 68.785081 | 183.927889 | |
| 7 | Male | 68.348516 | 167.971110 | |
| 8 | Male | 67.018950 | 175.929440 | |
| 9 | Male | 63.456494 | 156.399676 | |

Next steps:

[Generate code with data](#)[View recommended plots](#)

```
#checking the shape of the dataset
```

```
data.shape
```

```
(10000, 3)
```



```
#checking if there are null values in the dataset
```

```
data.isnull().sum()
```

```
Gender      0
Height      0
Weight      0
dtype: int64
```



```
#Getting a sample of the dataset
```

```
data.sample(10)
```

| | Gender | Height | Weight |  |
|------|--------|-----------|------------|---|
| 6211 | Female | 64.135856 | 132.991795 |  |
| 6253 | Female | 67.248224 | 172.713625 | |
| 1543 | Male | 68.469824 | 201.086014 | |
| 2686 | Male | 65.207347 | 151.875669 | |
| 9168 | Female | 64.453190 | 152.705133 | |
| 8533 | Female | 59.815911 | 130.957028 | |
| 2597 | Male | 68.957439 | 182.192513 | |
| 2008 | Male | 66.416414 | 193.601091 | |
| 2472 | Male | 70.074296 | 205.387910 | |
| 7655 | Female | 62.754392 | 131.292880 | |

```
#check bottom of the dataset
```

```
data.tail(10)
```

| | Gender | Height | Weight |  |
|------|--------|-----------|------------|---|
| 9990 | Female | 63.179498 | 141.266100 |  |
| 9991 | Female | 62.636675 | 102.853563 | |
| 9992 | Female | 62.077832 | 138.691680 | |
| 9993 | Female | 60.030434 | 97.687432 | |
| 9994 | Female | 59.098250 | 110.529686 | |
| 9995 | Female | 66.172652 | 136.777454 | |
| 9996 | Female | 67.067155 | 170.867906 | |
| 9997 | Female | 63.867992 | 128.475319 | |
| 9998 | Female | 69.034243 | 163.852461 | |
| 9999 | Female | 61.944246 | 113.649103 | |

```
column_to_drop = ['Gender']
```

```
#Drop column Gender
data = data.drop(column_to_drop, axis=1)
```

```
from sklearn.model_selection import train_test_split
```

```
#splitting the dataset
x_train, x_test, y_train, y_test = train_test_split(data.Height.values.reshape(-1, 1), data.Weight, random_state=11)
```

```
x_train.shape
(7500, 1)
```

```
y_train.shape
(7500,)
```

```
x_test.shape
(2500, 1)
```

```
from sklearn.linear_model import LinearRegression
```

```
#The linear regression model
linear_regression = LinearRegression()
```

```
#calling function fit so as to train the model
linear_regression.fit(X= x_train, y=y_train)
```

```
LinearRegression
LinearRegression()
```

```
#linear regression co-efficient (gradient of the slope)
linear_regression.coef_
array([7.71345358])
```

```
#linear regression y-intercept (where x and y meet)
linear_regression.intercept_
-350.390277050054
```

$$\text{equation} = y = mx + c$$

$$\text{weight} = 7.71345358 * \text{height} - 350.390277050054$$

```

#calling function predict to test the model
predicted_values = linear_regression.predict(x_test)

#re-assigning the expected values
expected_values = y_test

from sklearn.metrics import mean_absolute_error

#cheeking the Mean Abolute Error
print("MAE", mean_absolute_error(expected_values, predicted_values))

    MAE 9.585892229098492

#lambda function to calculate different weights
predict_weight = (lambda x: linear_regression.coef_ * x + linear_regression.intercept_)

#Weight predicted for a height of 70.1047862551571
predict_weight(70.1047862551571)

    array([190.35973736])

import matplotlib.pyplot as plt

weight = data['Weight']

weight

```

| | |
|------|------------|
| 0 | 241.893563 |
| 1 | 162.310473 |
| 2 | 212.740856 |
| 3 | 220.042470 |
| 4 | 206.349801 |
| ... | |
| 9995 | 136.777454 |
| 9996 | 170.867906 |
| 9997 | 128.475319 |
| 9998 | 163.852461 |
| 9999 | 113.649103 |

```

Name: Weight, Length: 10000, dtype: float64

height = data['Height']

height

```

| | |
|------|-----------|
| 0 | 73.847017 |
| 1 | 68.781904 |
| 2 | 74.110105 |
| 3 | 71.730978 |
| 4 | 69.881796 |
| ... | |
| 9995 | 66.172652 |
| 9996 | 67.067155 |
| 9997 | 63.867992 |
| 9998 | 69.034243 |
| 9999 | 61.944246 |

```

Name: Height, Length: 10000, dtype: float64

plt.figure(figsize=(20,20))

<Figure size 2000x2000 with 0 Axes>
<Figure size 2000x2000 with 0 Axes>

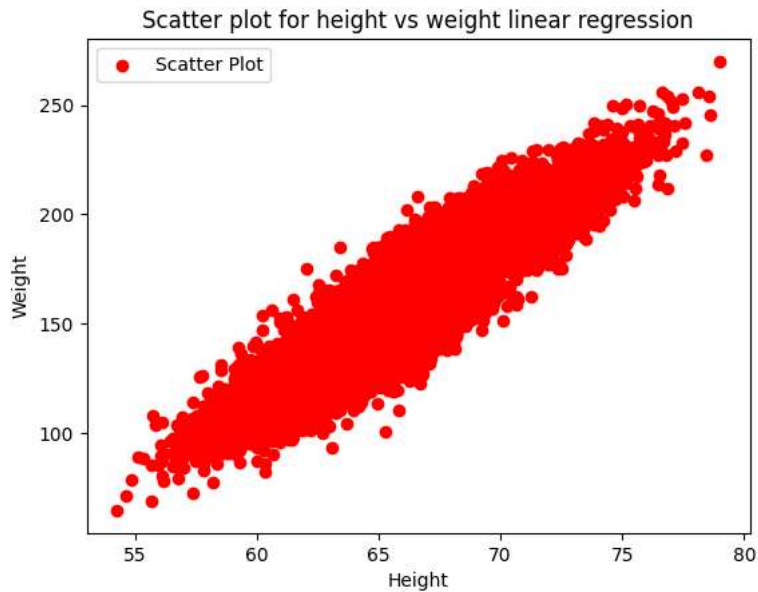
```

```
#scatter plot
#Attempt 1 at scatter plot
plt.scatter(height, weight, c='red', label='Scatter Plot')

#label for x-axis
plt.xlabel('Height')

#label for y-axis
plt.ylabel('Weight')

#Title for plot
plt.title("Scatter plot for height vs weight linear regression")
plt.legend()
plt.show()
```



Start coding or [generate](#) with AI.

```
import seaborn as sns
```

```
#seaborn scatterplot
#attempt number 2 of scatter plot
axes = sns.scatterplot(data=data, x='Height', y='Weight', hue='Height', palette='winter', legend=False)
```

```
x = np.array([min(data.Height.values), max(data.Height.values)])
y = predict(x)
line = plt.plot(x, y)
```

